

```
In [2]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
```

```
In [3]: 1 df=pd.read_csv(r"C:\Users\range\car_prediction_data.csv")
```

```
In [4]: 1 df.head()
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0

```
In [5]: 1 df.shape
```

```
(301, 9)
```

```
In [6]: 1 print(df['Seller_Type'].unique())
        2 print(df['Fuel_Type'].unique())
        3 print(df['Transmission'].unique())
        4 print(df['Owner'].unique())
```

```
['Dealer' 'Individual']
['Petrol' 'Diesel' 'CNG']
['Manual' 'Automatic']
[0 1 3]
```

```
In [7]: 1 df.isnull().sum()
```

```
Car_Name      0
Year          0
Selling_Price 0
Present_Price 0
Kms_Driven    0
Fuel_Type     0
Seller_Type   0
Transmission  0
Owner         0
dtype: int64
```

```
In [8]: 1 df.describe()
```

	Year	Selling_Price	Present_Price	Kms_Driven	Owner
count	301.000000	301.000000	301.000000	301.000000	301.000000
mean	2013.627907	4.661296	7.628472	36947.205980	0.043189
std	2.891554	5.082812	8.644115	38886.883882	0.247915
min	2003.000000	0.100000	0.320000	500.000000	0.000000
25%	2012.000000	0.900000	1.200000	15000.000000	0.000000
50%	2014.000000	3.600000	6.400000	32000.000000	0.000000
75%	2016.000000	6.000000	9.900000	48767.000000	0.000000
max	2018.000000	35.000000	92.600000	500000.000000	3.000000

```
In [9]: 1 final_dataset=df[['Year','Selling_Price','Present_Price','Kms_Driven','Fuel_Type','Seller_Type','Transm
```

```
In [10]: 1 final_dataset.head()
```

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0
1	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0
2	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0
3	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0
4	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0

```
In [11]: 1 final_dataset['Current_Year']=2021
```

```
In [12]: 1 final_dataset.head()
```

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Current_Year
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0	2021
1	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0	2021
2	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0	2021
3	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0	2021
4	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0	2021

```
In [13]: 1 final_dataset['no_year']=final_dataset['Current_Year']- final_dataset['Year']
```

```
In [14]: 1 final_dataset.head()
```

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Current_Year	no_year
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0	2021	7
1	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0	2021	8
2	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0	2021	4
3	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0	2021	10
4	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0	2021	7

```
In [15]: 1 final_dataset.drop(['Year'],axis=1,inplace=True)
```

```
In [16]: 1 final_dataset.head()
```

	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Current_Year	no_year
0	3.35	5.59	27000	Petrol	Dealer	Manual	0	2021	7
1	4.75	9.54	43000	Diesel	Dealer	Manual	0	2021	8
2	7.25	9.85	6900	Petrol	Dealer	Manual	0	2021	4
3	2.85	4.15	5200	Petrol	Dealer	Manual	0	2021	10
4	4.60	6.87	42450	Diesel	Dealer	Manual	0	2021	7

```
In [17]: 1 final_dataset=pd.get_dummies(final_dataset,drop_first=True)
```

```
In [18]: 1 final_dataset.head()
```

	Selling_Price	Present_Price	Kms_Driven	Owner	Current_Year	no_year	Fuel_Type_Diesel	Fuel_Type_Petrol	Seller_Type
0	3.35	5.59	27000	0	2021	7	False	True	False
1	4.75	9.54	43000	0	2021	8	True	False	False
2	7.25	9.85	6900	0	2021	4	False	True	False
3	2.85	4.15	5200	0	2021	10	False	True	False
4	4.60	6.87	42450	0	2021	7	True	False	False

```
In [19]: 1 final_dataset=final_dataset.drop(['Current_Year'],axis=1)
```

```
In [20]: 1 final_dataset.head()
```

	Selling_Price	Present_Price	Kms_Driven	Owner	no_year	Fuel_Type_Diesel	Fuel_Type_Petrol	Seller_Type_Individual	Transmission
0	3.35	5.59	27000	0	7	False	True	False	True
1	4.75	9.54	43000	0	8	True	False	False	True
2	7.25	9.85	6900	0	4	False	True	False	True
3	2.85	4.15	5200	0	10	False	True	False	True
4	4.60	6.87	42450	0	7	True	False	False	True

```
In [21]: 1 final_dataset.corr()
```


	Selling_Price	Present_Price	Kms_Driven	Owner	no_year	Fuel_Type_Diesel	Fuel_Type_Petrol	Seller_Type_Individual	Transmission
Selling_Price	1.000000	0.878983	0.029187	-0.088344	-0.236141	0.552339	-0.540571	-0.550724	-0.367128
Present_Price	0.878983	1.000000	0.203647	0.008057	0.047584	0.473306	-0.465244	-0.512030	-0.348715
Kms_Driven	0.029187	0.203647	1.000000	0.089216	0.524342	0.172515	-0.172874	-0.101419	-0.162510
Owner	-0.088344	0.008057	0.089216	1.000000	0.182104	-0.053469	0.055687	0.124269	-0.050316
no_year	-0.236141	0.047584	0.524342	0.182104	1.000000	-0.064315	0.059959	0.039896	-0.000394
Fuel_Type_Diesel	0.552339	0.473306	0.172515	-0.053469	-0.064315	1.000000	-0.979648	-0.350467	-0.098643
Fuel_Type_Petrol	-0.540571	-0.465244	-0.172874	0.055687	0.059959	-0.979648	1.000000	0.358321	0.091013
Seller_Type_Individual	-0.550724	-0.512030	-0.101419	0.124269	0.039896	-0.350467	0.358321	1.000000	0.063245
Transmission	-0.367128	-0.348715	-0.162510	-0.050316	-0.000394	-0.098643	0.091013	0.063245	1.000000


```
In [22]: 1 import seaborn as sns
```

```
In [23]: 1 sns.pairplot(final_dataset)
```

```
<__array_function__ internals>:200: RuntimeWarning: Converting input from bool to <class 'numpy.uint8'> for compatibility.  
<__array_function__ internals>:200: RuntimeWarning: Converting input from bool to <class 'numpy.uint8'> for compatibility.  
<__array_function__ internals>:200: RuntimeWarning: Converting input from bool to <class 'numpy.uint8'> for compatibility.  
<__array_function__ internals>:200: RuntimeWarning: Converting input from bool to <class 'numpy.uint8'> for compatibility.  
C:\Users\range\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight  
    self._figure.tight_layout(*args, **kwargs)
```

```
<seaborn.axisgrid.PairGrid at 0x1c9456e1e50>
```





```
In [25]: 1 X=final_dataset.iloc[:,1:]
          2 y=final_dataset.iloc[:,0]
```

```
In [26]: 1 X['Owner'].unique()

          array([0, 1, 3], dtype=int64)
```

```
In [27]: 1 X.head()
```

	Present_Price	Kms_Driven	Owner	no_year	Fuel_Type_Diesel	Fuel_Type_Petrol	Seller_Type_Individual	Transmission_Ma
0	5.59	27000	0	7	False	True	False	True
1	9.54	43000	0	8	True	False	False	True
2	9.85	6900	0	4	False	True	False	True
3	4.15	5200	0	10	False	True	False	True
4	6.87	42450	0	7	True	False	False	True

```
In [28]: 1 y.head()

0    3.35
1    4.75
2    7.25
3    2.85
4    4.60
Name: Selling_Price, dtype: float64
```

Feature Importance

```
In [29]: 1 from sklearn.ensemble import ExtraTreesRegressor
2 model = ExtraTreesRegressor()
3 model.fit(X,y)
```


▼ ExtraTreesRegressor

ExtraTreesRegressor()

```
In [30]: 1 print(model.feature_importances_)

[0.37028483 0.03825943 0.00078987 0.07634868 0.24050773 0.00745442
 0.12198921 0.14436582]
```

```
In [31]: 1 feat_importances = pd.Series(model.feature_importances_, index=X.columns)
          2 feat_importances.nlargest(5).plot(kind='barh')
          3 plt.show()
```



```
In [32]: 1 from sklearn.model_selection import train_test_split
          2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
```

Using Random Forest Regressor

In [33]:	1	<code>from sklearn.ensemble import RandomForestRegressor</code>
In [34]:	1	<code>regressor=RandomForestRegressor()</code>
In [35]:	1	<code>from sklearn.model_selection import RandomizedSearchCV</code>
In [36]:	1 2 3 4 5 6 7 8	<pre><i>#Randomized Search CV</i> n_estimators = [int(x) for x in np.linspace(start = 100, stop = 1200, num = 12)] max_features = ['auto', 'sqrt'] max_depth = [int(x) for x in np.linspace(5, 30, num = 6)] min_samples_split = [2, 5, 10, 15, 100] min_samples_leaf = [1, 2, 5, 10]</pre>
In [37]:	1 2 3 4 5 6 7	<pre>random_grid = {'n_estimators': n_estimators, 'max_features': max_features, 'max_depth': max_depth, 'min_samples_split': min_samples_split, 'min_samples_leaf': min_samples_leaf} print(random_grid)</pre>
		<pre>{'n_estimators': [100, 200, 300, 400, 500, 600, 700, 800, 900, 1000, 1100, 1200], 'max_features': ['auto', 'sqrt'], 'max_depth': [5, 10, 15, 20, 25, 30], 'min_samples_split': [2, 5, 10, 15, 100], 'min_samples_leaf': [1, 2, 5, 10]}</pre>
In [38]:	1	<code>rf = RandomForestRegressor()</code>

```
In [39]: 1 # Use the random grid to search for best hyperparameters
        2 rf=RandomizedSearchCV(estimator = rf, param_distributions = random_grid,scoring='neg_mean_squared_error
```

```
In [40]: 1 rf.fit(X_train,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[CV] END max_depth=10, max_features=sqrt, min_samples_leaf=5, min_samples_split=5, n_estimators=900; total time= 1.5s
[CV] END max_depth=10, max_features=sqrt, min_samples_leaf=5, min_samples_split=5, n_estimators=900; total time= 1.5s
[CV] END max_depth=10, max_features=sqrt, min_samples_leaf=5, min_samples_split=5, n_estimators=900; total time= 1.4s
[CV] END max_depth=10, max_features=sqrt, min_samples_leaf=5, min_samples_split=5, n_estimators=900; total time= 1.3s
[CV] END max_depth=10, max_features=sqrt, min_samples_leaf=5, min_samples_split=5, n_estimators=900; total time= 1.4s
[CV] END max_depth=15, max_features=sqrt, min_samples_leaf=2, min_samples_split=10, n_estimators=1100; total time= 1.7s
[CV] END max_depth=15, max_features=sqrt, min_samples_leaf=2, min_samples_split=10, n_estimators=1100; total time= 1.7s
[CV] END max_depth=15, max_features=sqrt, min_samples_leaf=2, min_samples_split=10, n_estimators=1100; total time= 1.6s
[CV] END max_depth=15, max_features=sqrt, min_samples_leaf=2, min_samples_split=10, n_estimators=1100; total time= 1.8s
[CV] END max_depth=15, max_features=sqrt, min_samples_leaf=2, min_samples_split=10, n_estimators=1100; total time= 1.9s
```

C:\Users\range\anaconda3\Lib\site-packages\sklearn\ensemble_forest.py:413: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

warn(

```
[CV] END max_depth=15, max_features=auto, min_samples_leaf=5, min_samples_split=100, n_estimators=300; total time= 0.5s
```

C:\Users\range\anaconda3\Lib\site-packages\sklearn\ensemble_forest.py:413: FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the def

```
In [41]: 1 rf.best_params_
```

```
{'n_estimators': 1000,
 'min_samples_split': 2,
 'min_samples_leaf': 1,
 'max_features': 'sqrt',
 'max_depth': 25}
```

```
In [42]: 1 rf.best_score_
```

```
-3.9930132501533486
```

```
In [43]: 1 predictions=rf.predict(X_test)
```

```
In [44]: 1 sns.distplot(y_test-predictions)
```

```
C:\Users\range\AppData\Local\Temp\ipykernel_1276\2131792714.py:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

```
Please adapt your code to use either `displot` (a figure-level function with  
similar flexibility) or `histplot` (an axes-level function for histograms).
```

```
For a guide to updating your code to use the new functions, please see  
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751 (https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751)
```


```
sns.distplot(y_test-predictions)
```

```
<Axes: xlabel='Selling_Price', ylabel='Density'>
```




```
In [45]: 1 plt.scatter(y_test,predictions)

<matplotlib.collections.PathCollection at 0x1c94f695390>
```



```
In [46]: 1 from sklearn import metrics
```

```
In [47]: 1 print('MAE:', metrics.mean_absolute_error(y_test, predictions))
          2 print('MSE:', metrics.mean_squared_error(y_test, predictions))
          3 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

```
MAE: 0.887862417582416
MSE: 3.9217544933780215
RMSE: 1.9803420142435046
```

Using XGBoost Regressor

```
In [48]: 1 import xgboost as xgb
          2 from scipy.stats import uniform, randint
```

```
In [49]: 1 xgb_model = xgb.XGBRegressor(objective="reg:linear", random_state=42)
```

```
In [50]: 1 params = {
          2     "gamma": uniform(0, 0.5),
          3     "learning_rate": uniform(0.03, 0.3), # default 0.1
          4     "max_depth": randint(2, 6), # default 3
          5     "n_estimators": randint(100, 150), # default 100
          6     "subsample": uniform(0.6, 0.4)
          7 }
```

```
In [51]: 1 xgb = RandomizedSearchCV(estimator = xgb_model, param_distributions = params, scoring='neg_mean_squared_
```

```
In [52]: 1 xgb.fit(X_train,y_train)
```

```
Fitting 5 folds for each of 10 candidates, totalling 50 fits
```

```
[CV] END gamma=0.18727005942368125, learning_rate=0.3152142919229748, max_depth=4, n_estimators=107, subsample=0.8394633936788146; total time= 0.0s
```

```
[CV] END gamma=0.18727005942368125, learning_rate=0.3152142919229748, max_depth=4, n_estimators=107, subsample=0.8394633936788146; total time= 0.0s
```

```
C:\Users\range\anaconda3\Lib\site-packages\xgboost\core.py:160: UserWarning: [12:18:13] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07f6e447eee219473-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:209: reg:linear is now deprecated in favor of reg:squarederror.
```

```
warnings.warn(msg, UserWarning)
```

```
C:\Users\range\anaconda3\Lib\site-packages\xgboost\core.py:160: UserWarning: [12:18:13] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07f6e447eee219473-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:209: reg:linear is now deprecated in favor of reg:squarederror.
```

```
warnings.warn(msg, UserWarning)
```

```
C:\Users\range\anaconda3\Lib\site-packages\xgboost\core.py:160: UserWarning: [12:18:13] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07f6e447eee219473-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:209: reg:linear is now deprecated in favor of reg:squarederror.
```

```
warnings.warn(msg, UserWarning)
```

```
[CV] END gamma=0.18727005942368125, learning_rate=0.3152142919229748, max_depth=4, n_estimators=107, subsample=0.8394633936788146; total time= 0.0s
```

```
[CV] END gamma=0.18727005942368125, learning_rate=0.3152142919229748, max_depth=4, n_estimators=107, subsample=0.8394633936788146; total t
```

```
In [53]: 1 xgb.best_score_
```

```
-1.9114686351785255
```

```
In [54]: 1 xgb.best_params_  
  
{'gamma': 0.3005575058716044,  
 'learning_rate': 0.24242177333881365,  
 'max_depth': 3,  
 'n_estimators': 101,  
 'subsample': 0.8887995089067299}
```

```
In [55]: 1 predictions=xgb.predict(X_test)
```

```
In [56]: 1 sns.distplot(y_test-predictions)
```

```
C:\Users\range\AppData\Local\Temp\ipykernel_1276\2131792714.py:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

```
Please adapt your code to use either `displot` (a figure-level function with  
similar flexibility) or `histplot` (an axes-level function for histograms).
```

```
For a guide to updating your code to use the new functions, please see  
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751 (https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751)
```


```
sns.distplot(y_test-predictions)
```

```
<Axes: xlabel='Selling_Price', ylabel='Density'>
```



```
In [57]: 1 plt.scatter(y_test,predictions)

<matplotlib.collections.PathCollection at 0x1c95631e9d0>
```



```
In [58]: 1 print('MAE:', metrics.mean_absolute_error(y_test, predictions))
          2 print('MSE:', metrics.mean_squared_error(y_test, predictions))
          3 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

MAE: 0.7972087596921802

MSE: 2.6811867849257176

RMSE: 1.637432986392334

Using Catboost Regressor

```
In [62]: 1 from catboost import CatBoostRegressor
```



```
In [63]: 1 !pip install catboost
```

```
Requirement already satisfied: catboost in c:\users\range\anaconda3\lib\site-packages (1.2.3)
Requirement already satisfied: graphviz in c:\users\range\anaconda3\lib\site-packages (from catboost) (0.20.3)
Requirement already satisfied: matplotlib in c:\users\range\anaconda3\lib\site-packages (from catboost) (3.7.2)
Requirement already satisfied: numpy>=1.16.0 in c:\users\range\anaconda3\lib\site-packages (from catboost) (1.24.3)
Requirement already satisfied: pandas>=0.24 in c:\users\range\anaconda3\lib\site-packages (from catboost) (2.0.3)
Requirement already satisfied: scipy in c:\users\range\anaconda3\lib\site-packages (from catboost) (1.11.1)
Requirement already satisfied: plotly in c:\users\range\anaconda3\lib\site-packages (from catboost) (5.9.0)
Requirement already satisfied: six in c:\users\range\anaconda3\lib\site-packages (from catboost) (1.16.0)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\range\anaconda3\lib\site-packages (from pandas>=0.24->catboost) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in c:\users\range\anaconda3\lib\site-packages (from pandas>=0.24->catboost) (2023.3.post1)
Requirement already satisfied: tzdata>=2022.1 in c:\users\range\anaconda3\lib\site-packages (from pandas>=0.24->catboost) (2023.3)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\range\anaconda3\lib\site-packages (from matplotlib->catboost) (1.0.5)
Requirement already satisfied: cycler>=0.10 in c:\users\range\anaconda3\lib\site-packages (from matplotlib->catboost) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\range\anaconda3\lib\site-packages (from matplotlib->catboost) (4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\range\anaconda3\lib\site-packages (from matplotlib->catboost) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\range\anaconda3\lib\site-packages (from matplotlib->catboost) (23.1)
Requirement already satisfied: pillow>=6.2.0 in c:\users\range\anaconda3\lib\site-packages (from matplotlib->catboost) (9.4.0)
Requirement already satisfied: pyparsing<3.1,>=2.3.1 in c:\users\range\anaconda3\lib\site-packages (from matplotlib->catboost) (3.0.9)
Requirement already satisfied: tenacity>=6.2.0 in c:\users\range\anaconda3\lib\site-packages (from plotly->catboost) (8.2.2)
```

```
[notice] A new release of pip is available: 23.3.2 -> 24.0
```

```
[notice] To update, run: python.exe -m pip install --upgrade pip
```

```
In [64]: 1 cb=CatBoostRegressor()
```

```
In [65]: 1 grid = {'learning_rate': [0.03, 0.1],
2             'depth': [4, 6, 10],
3             'l2_leaf_reg': [1, 3, 5, 7, 9]}
```

```
In [66]: 1 cb = RandomizedSearchCV(estimator = cb, param_distributions = grid, scoring='neg_mean_squared_error', n_
```

```
In [67]: 1 cb.fit(X_train,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

0:	learn: 4.9398622	total: 136ms	remaining: 2m 15s
1:	learn: 4.7805986	total: 143ms	remaining: 1m 11s
2:	learn: 4.5730251	total: 146ms	remaining: 48.7s
3:	learn: 4.4161829	total: 149ms	remaining: 37s
4:	learn: 4.2856056	total: 159ms	remaining: 31.7s
5:	learn: 4.1268798	total: 161ms	remaining: 26.6s
6:	learn: 3.9895888	total: 164ms	remaining: 23.3s
7:	learn: 3.8391546	total: 169ms	remaining: 20.9s
8:	learn: 3.7310291	total: 182ms	remaining: 20s
9:	learn: 3.6173828	total: 183ms	remaining: 18.1s
10:	learn: 3.4712492	total: 184ms	remaining: 16.6s
11:	learn: 3.3483253	total: 187ms	remaining: 15.4s
12:	learn: 3.2450271	total: 194ms	remaining: 14.8s
13:	learn: 3.1400164	total: 196ms	remaining: 13.8s
14:	learn: 3.0260343	total: 198ms	remaining: 13s
15:	learn: 2.9346526	total: 199ms	remaining: 12.3s
16:	learn: 2.8376114	total: 211ms	remaining: 12.2s
17:	learn: 2.7572800	total: 213ms	remaining: 11.6s
18:	learn: 2.6724960	total: 224ms	remaining: 11.6s
19:	learn: 2.6150300	total: 234ms	remaining: 11.5s
20:	learn: 2.5479388	total: 237ms	remaining: 11s
21:	learn: 2.4920401	total: 240ms	remaining: 10.7s
22:	learn: 2.4787222	total: 241ms	remaining: 10.2s

```
In [68]: 1 cb.best_score_
```

-3.2179271237932716

```
In [69]: 1 cb.best_params_
```

{'learning_rate': 0.1, 'l2_leaf_reg': 9, 'depth': 4}

```
In [70]: 1 predictions=cb.predict(X_test)
```

```
In [71]: 1 sns.distplot(y_test-predictions)
```

C:\Users\range\AppData\Local\Temp\ipykernel_1276\2131792714.py:1: UserWarning:


``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)


```
sns.distplot(y_test-predictions)
```

```
<Axes: xlabel='Selling_Price', ylabel='Density'>
```



```
In [72]: 1 plt.scatter(y_test,predictions)

<matplotlib.collections.PathCollection at 0x1c958859590>
```



```
In [73]: 1 print('MAE:', metrics.mean_absolute_error(y_test, predictions))
          2 print('MSE:', metrics.mean_squared_error(y_test, predictions))
          3 print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

MAE: 0.8568669682476366

MSE: 3.4731227057809524

RMSE: 1.8636315906801302

```
In [ ]: 1
```